

NLP and Generative AI

Spikeup.ai

16 August 2023

Nice to meet you!

Your facilitators for this training



Anshul Aggarwal

Consultant - Artificial Intelligence

~1 year at Deloitte
Background in
Computer Science / AI



Alex Jang

Analyst - Artificial Intelligence

~1 year at Deloitte
Background in
Computer Science / AI

1 NLP tasks and feature representation

1. NLP preprocessing tasks
2. NLP feature representation
3. Neural networks



NLP has been split up in so-called “tasks”

Each task has specific use cases, models and evaluation metrics



Information Retrieval

Ranking a list of documents or search results in response to an input query



Reading comprehension

Answering questions on a given text passage



Sentiment Analysis

Classifying the polarity or “happiness” of a given text



Summarization

Producing a shorter version of one or several documents that preserves most of the input’s meaning



Machine Translation

Translating a sentence in a source language to a different target language




Topic modelling

Discovering the abstract “topics” that occur in a collection of documents



Natural language inference

Determining whether a “hypothesis” is true (entailment), false (contradiction), or undetermined (neutral) given a “premise”





Relation prediction

Recognizing a named relation between two named semantic entities



Sentiment analysis

Given a typically short piece of text, classify the emotional tone of the message: positive, negative, or neutral

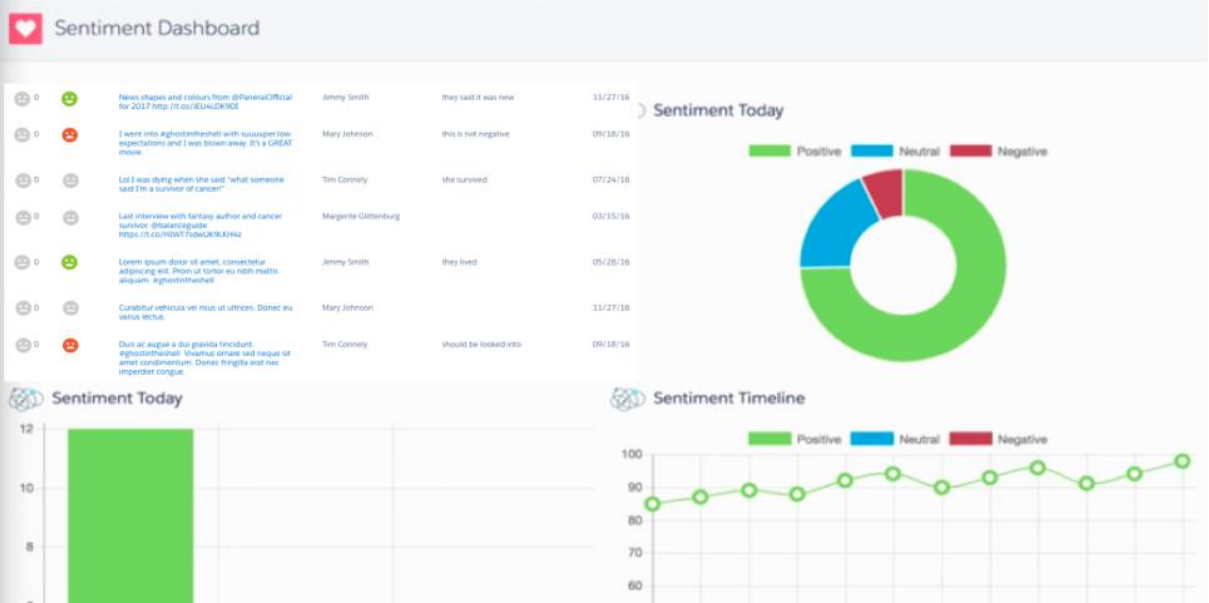
See what guests loved the most:

**Mandy**
 Germany

"The design was so cute and love the tiles in the bathroom "

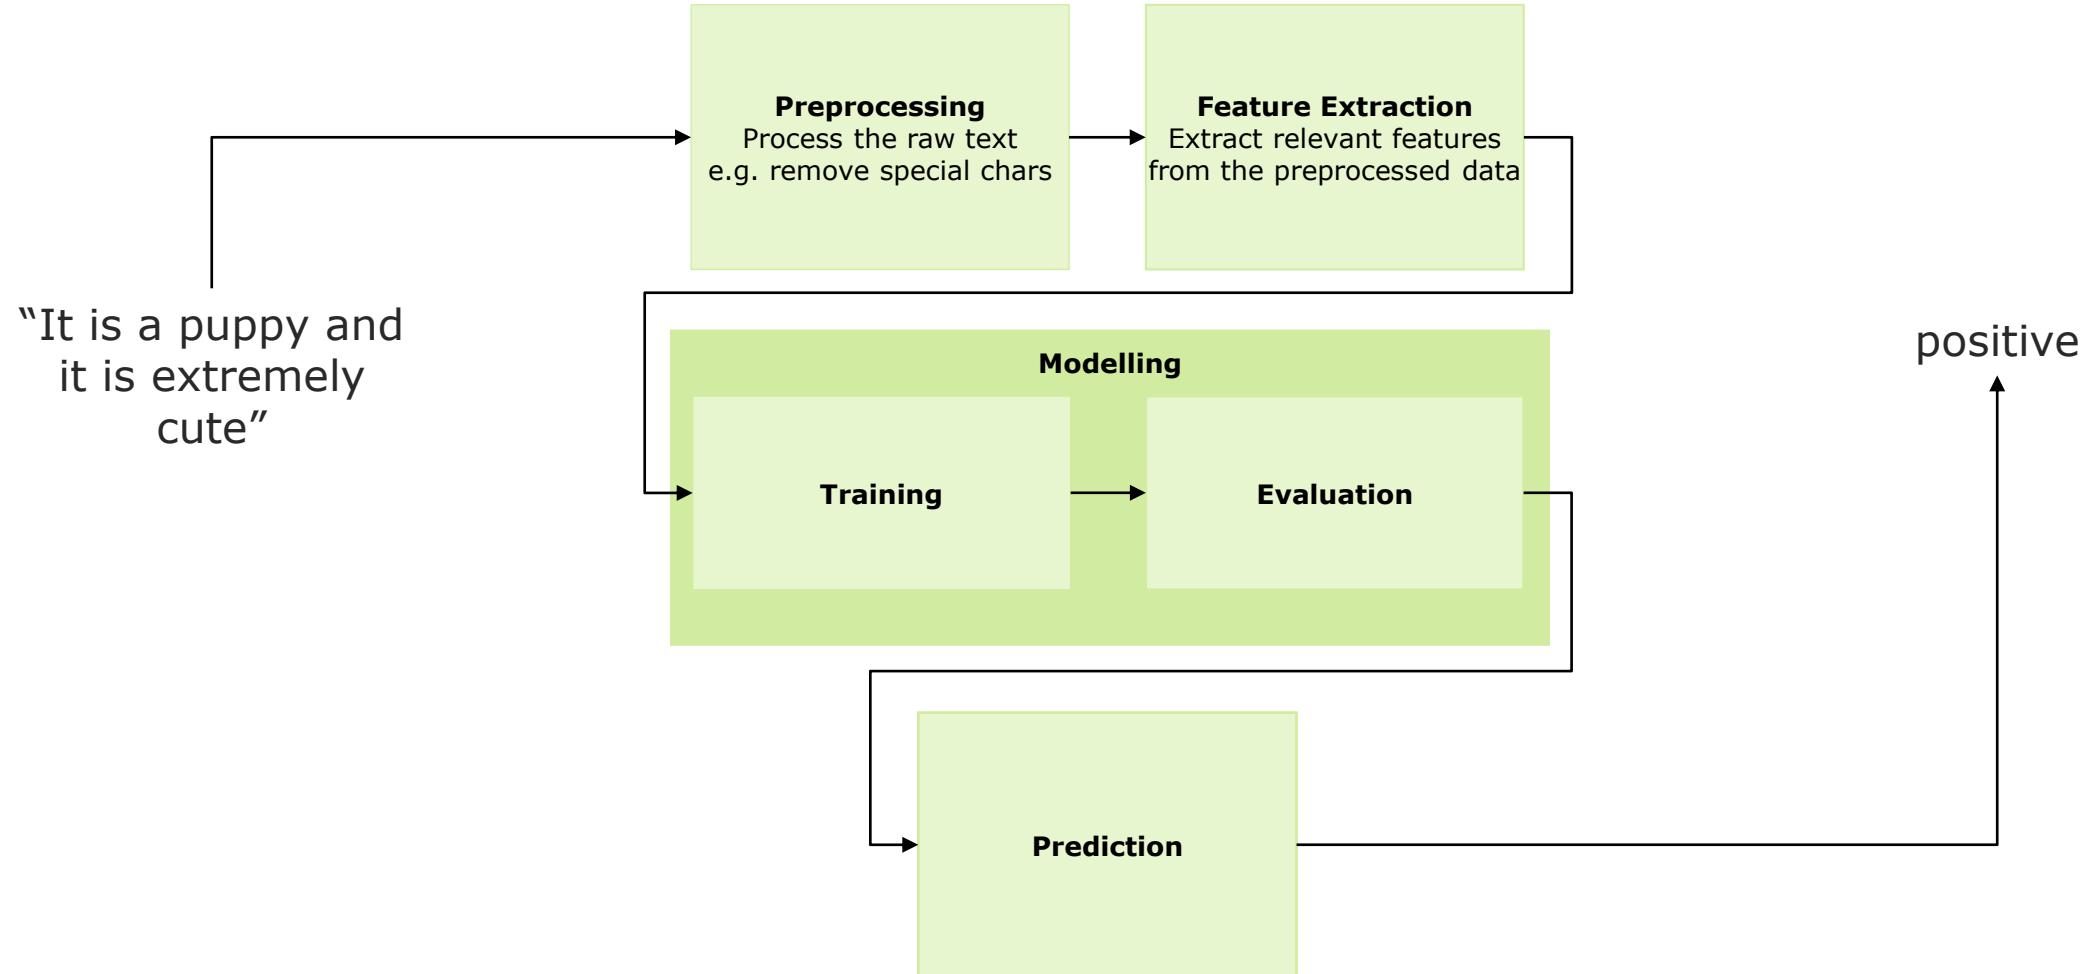
**Marieke**
 Netherlands

"I really liked the atmosphere. Nicely decorated. Great location."



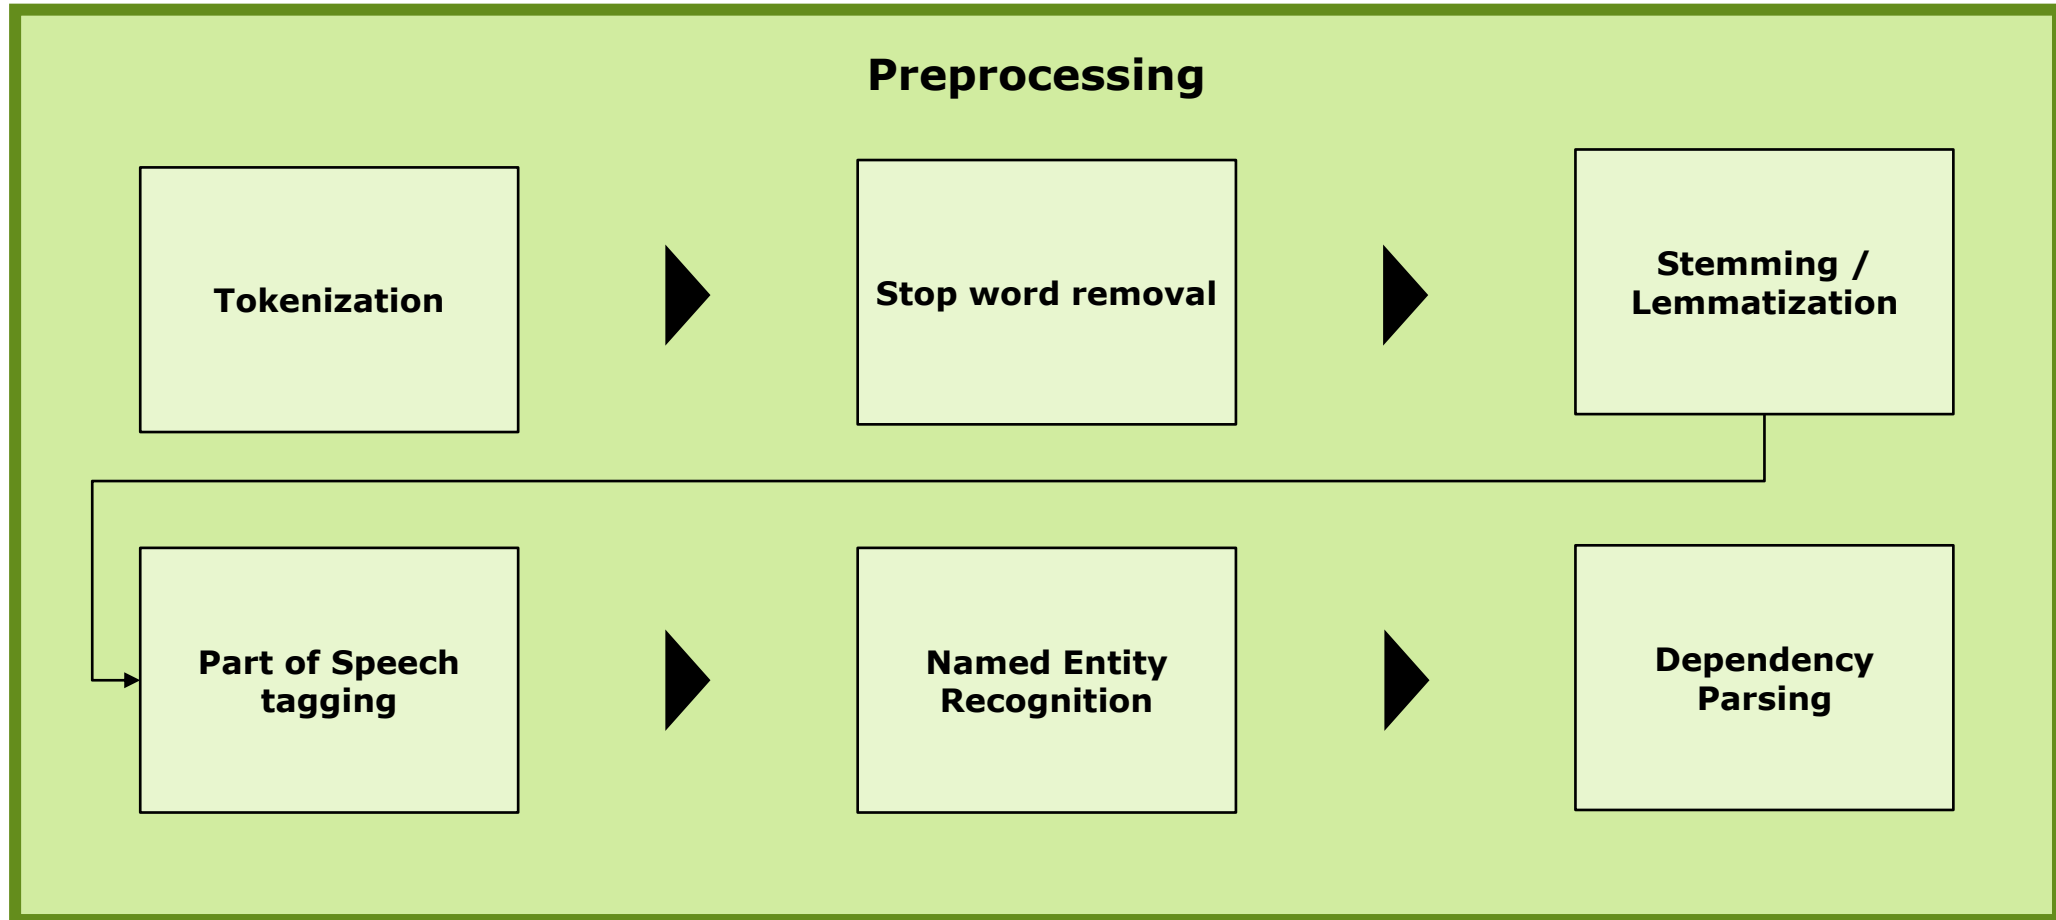
Sentiment analysis

Given a typically short piece of text, classify the emotional tone of the message: positive, negative, or neutral



NLP preprocessing pipeline

What NLP preprocessing would you apply to get the sentiment per sentence?



NLP Preprocessing tasks

Several basic NLP subtasks exist to transform raw data to information



Tokenization

Replace a sequence of characters with a sequence of tokens (words and punctuation)




Stop word removal

Remove ubiquitous words from a sentence



Stemming

Cutting "inflected" words to their root forms by removing suffixes



Lemmatization

Group a so-called "inflected" form of a word to its dictionary form ("lemma")




Part of Speech (PoS) tagging

Assigning parts of speech to individual words in a sentence



Named Entity Recognition (NER)

Identifying important words in a sentence



Dependency parsing

Identify how words relate to each other.

NLP subtasks and pipeline

Example outputs



Tokenization

"I like London,"




Stop word removal

"The city is large"



Stemming

"prettiest"



Lemmatization

"is"




Part of Speech (PoS) tagging

"Large city"



Named Entity Recognition (NER)

"London"



Dependency parsing

"London is the capital"

NLP subtasks and pipeline

Example outputs



Tokenization

"I like London,"

["I", "like", "London", ",", ""]



Stop word removal

"The city is large"


["City", "is", "large"]



Stemming

"prettiest"

"pretti"



Lemmatization

"is"

"be"



Part of Speech (PoS) tagging

"Large city"


{"Large": "adjective", "city": "noun"}



Named Entity Recognition (NER)


"London"

{"London": "geographical entity"}



Dependency parsing

"London is the capital"



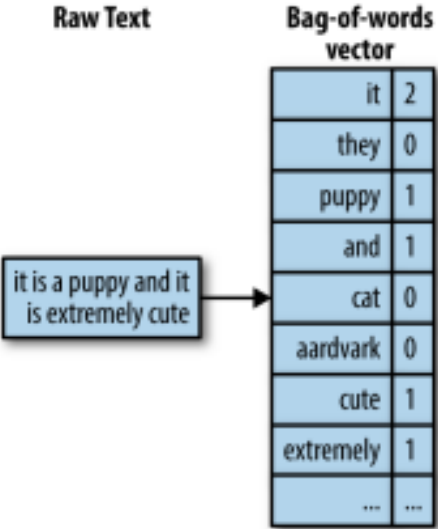
NLP and features

Encoding natural language into numerical values for computation


Bag of Words (BOW): Represent using presence/count of words in a piece of text with the idea that documents are similar if they contain the same words

Bag of n-grams: Presence/count of bigrams/trigrams/n-grams; more informative since they capture more context around each word

Vector representation: These gives us a vector representation of natural language text that we can perform computations with.



Bigram	
it is	1
is a	1
...	...
extremely cute	1
not cute	0



Embeddings

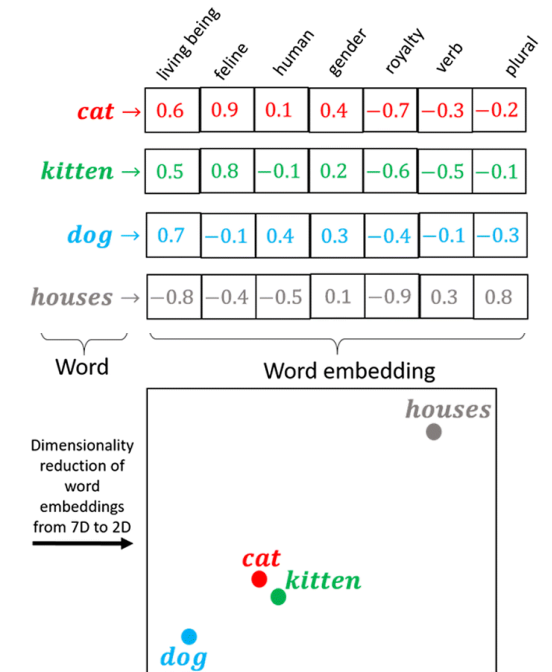
A special type of feature representation

Embeddings: Latent representations that preserve semantics.

- Compressed feature representation of data that contain semantic information
- Features along the dimensions are directly interpretable.
- Convert **words or sentences** into **vectors that preserves its meanings and structure**.
- Does not create deeper features. They are **generated using simpler architectures**.
- Can be used to compare pairs of words or sentences – do they mean something similar?
 - You can use cosine similarity!
- Are used in all LLMs.

Embeddings (e.g. word2vec)


let NN identify patterns




Neural networks

Artificial neurons that mimic the human brain


- **Neural network:** neurons organised in layers; information flows through the layers to give final output



- **Encoders:** produces a vector representation
encodes relevant information from the input

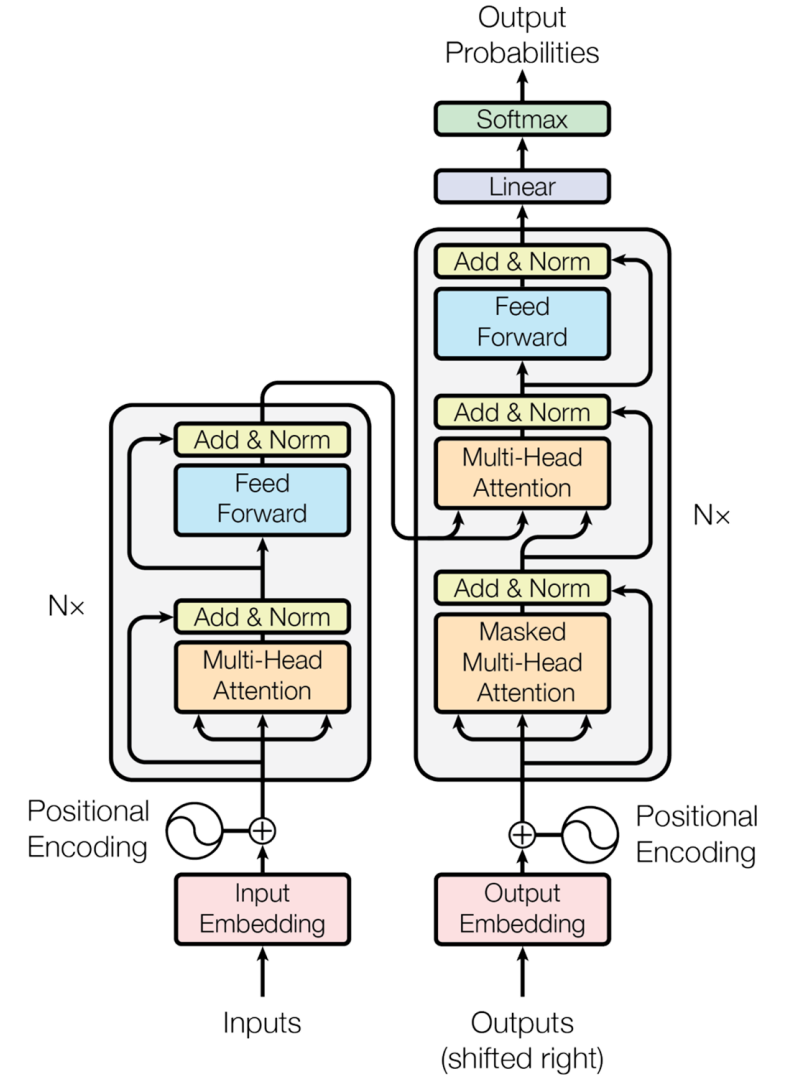


- **Decoders:** produces an output result for the task
by **decoding** the information given



The Transformer changed the (NLP) world

Source: A. Vaswani et al., "Attention is all you need", June 2017



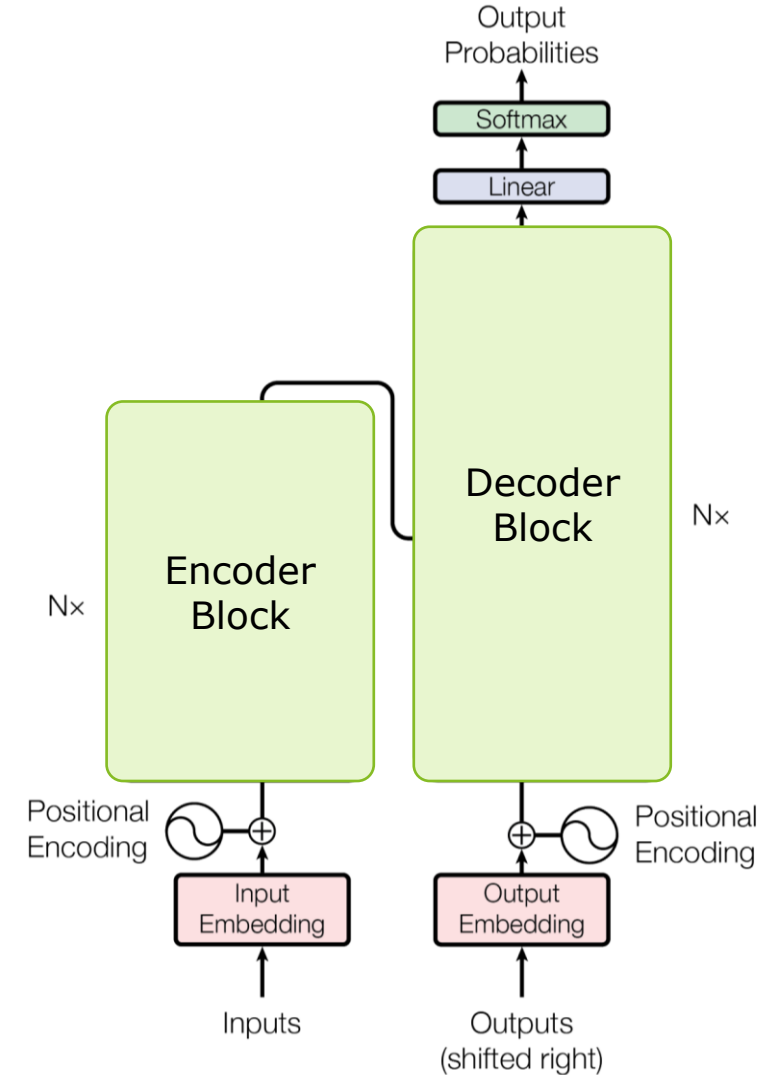

The fundamental structure is familiar

It builds up on the encoder-decoder architecture

Latent representation: Compressed feature representation of data. Called “latent” (*hidden*) as it cannot be directly observed but is inferred from the input.

Encoder: A neural network component that converts input data, into a latent representation, usually a vector of fixed size.

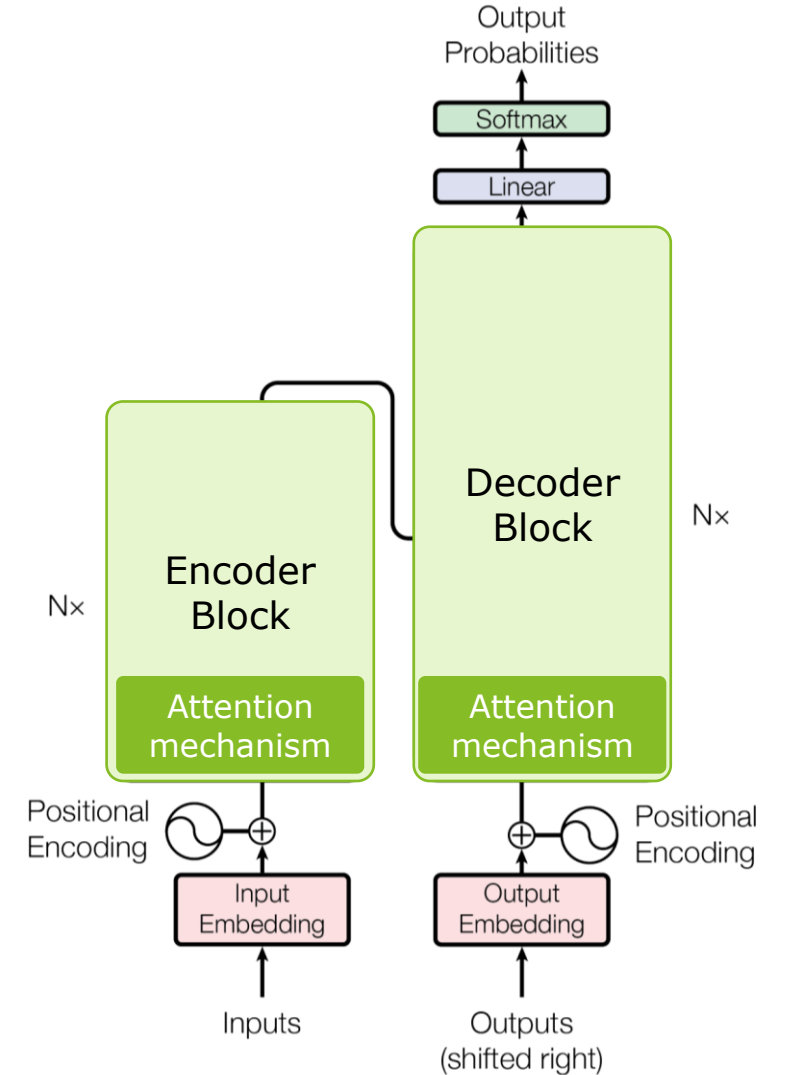
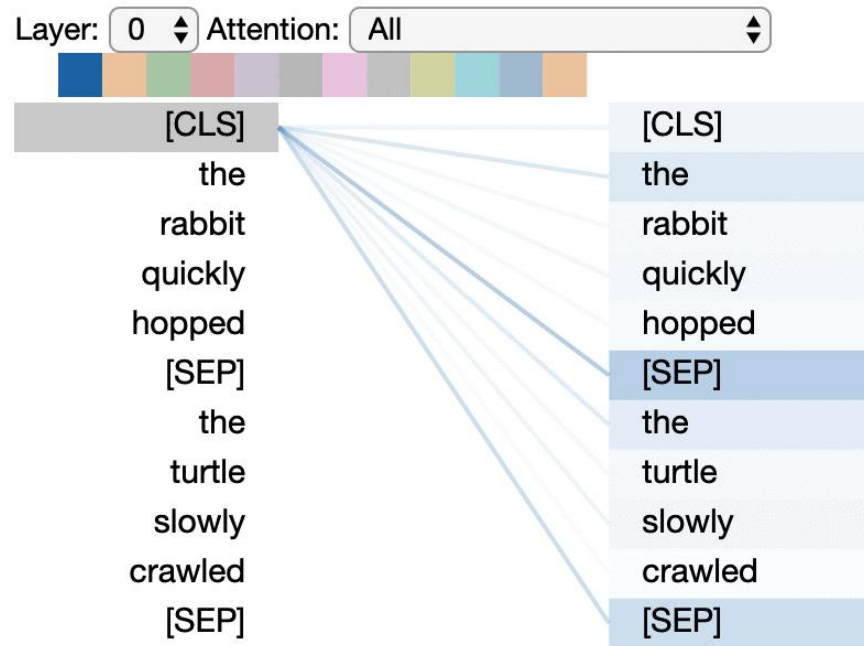
Decoder: The opposite of an Encoder, it converts a latent representation into the desired output form.



Attention is all you need

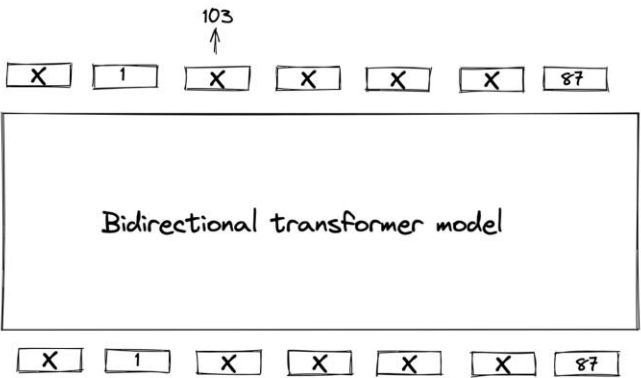
Understanding context in a better way

Attention mechanism: Increase the importance of some parts of the input (*pay attention to these parts*), while diminishing the importance of others for a given token. What parts to focus on and by how much, is learnt during the training process.

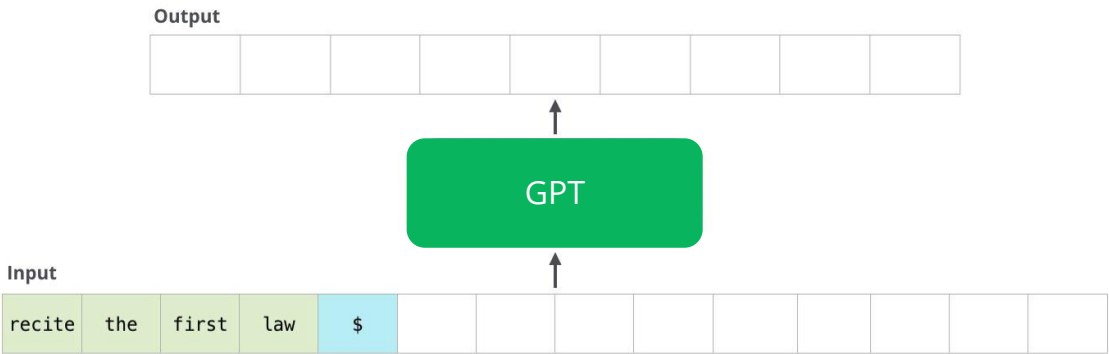


Comparing two different families of LLMs

	BERT – Bidirectional Encoder Representations from Transformers	GPT – Generative Pretrained Transformer
CONTEXT DIRECTION	Bidirectional: Considers both the left and right contexts to make predictions.	Autoregressive: Considers only the context on the left to make predictions
TASKS	Suitable for Natural Language Understanding tasks.	Suitable for Natural Language Understanding and text generation tasks.
ARCHITECTURE	Contains only the encoder blocks of the transformer architecture	Contains only the decoder blocks of the transformer architecture
FINE-TUNING	Necessary for all task-specific operations.	Only used when memorisation of private content or simulation of writing style/tone of voice is required.
MODEL DEPLOYMENT	A task specific model is deployed for applications	A general purpose model can be deployed , using prompts for defining tasks
DELOITTE USE CASES	Classify customer feedback at ING, using RoBERTa	Make internal Knowledge easily accessible using GPT-3 and Search, “QueryGenie” for Shell



Source: B. Christiaens., Potential applications of discrete diffusion models”, August 2022

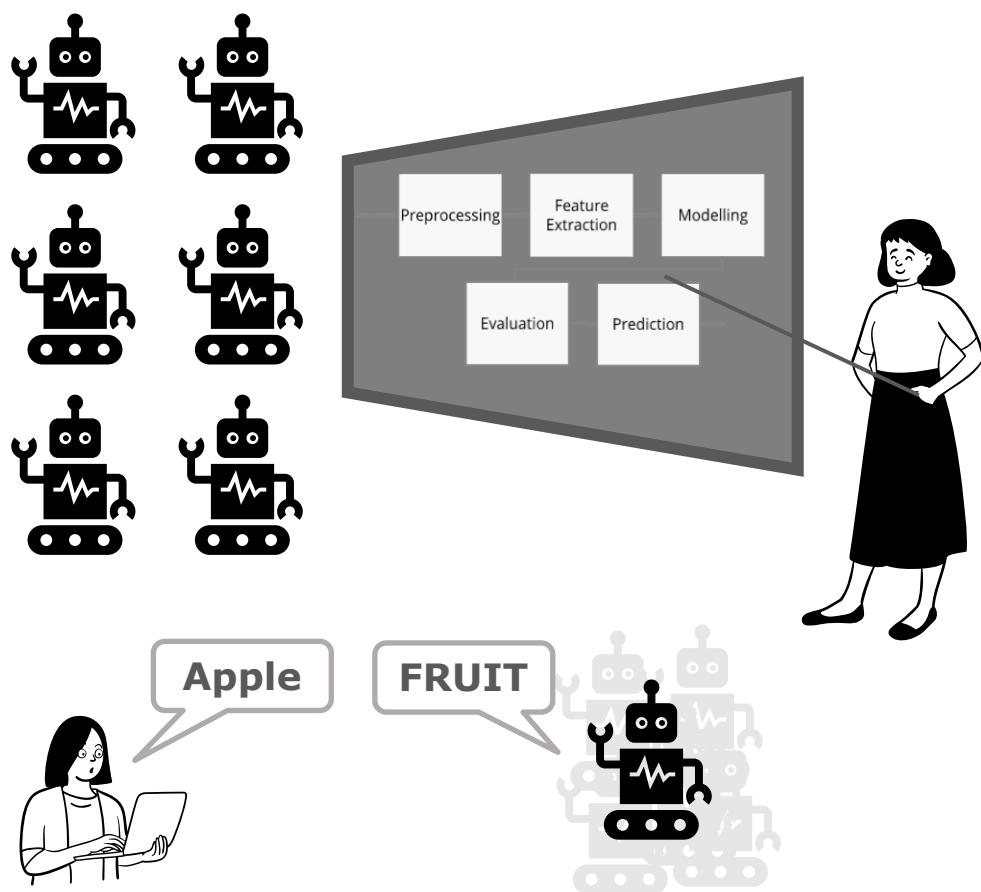


Source: J. Alammari, “The Illustrated GPT-2 (Visualizing Transformer Language Models)”, August 2019


Generative LLMs can take over the entire NLP flow

They can be deployed easily end-to-end, simplifying NLP application development

What NLP looked like before Generative LLMs – a lot of training for every task



What NLP looks like now – just ask!



Prompt engineering

Improving your prompts


Never forget these tips that will help you take your Generative LLM results to a higher level



Be **Concise** and **Specific**



Include your **own Knowledge**



Identify the **Goal**



Create a **Draft Output**



Prime the AI


The Power of Prompting

Be clear and specific to what you expect out of it. Do not provide misleading information that can confuse the model

Hey GPT-3, would you be kind enough to tell me when was the last time the out-of-solar-system spacecraft (I really forgot its name, sorry, could you also let me know what was it?) sent out a message to us? Also it was sent by the Japanese, correct? Thank you!

Submit Ctrl Enter


Submit ↶ ↷ 68



Name the spacecrafts have reached interstellar space? Who sent them, and are they still in contact with earth?


Submit Ctrl Enter

Submit ↶ ↷ 22



The Power of Prompting

Set the scene, or provide some examples

Pretend that you are Siri, a virtual assistant, for users based in the Netherlands. A valid user command starts with the exact words "Hey Siri". For commands that do not start with "Hey Siri", ignore the command and respond with the word "crickets". Pretend that you can execute the valid commands and respond with a dummy message. Create a response to the following user commands. 

User 1: Hey Siri, can you play a rap song?

User 2: Play Beethoven's Fifth Symphony.

User 3: Hey Siri, how many steps did I walk today?


User 4: Hey Siri, how's the weather forecast for tomorrow?

User 5: Hey Siri, set an alarm for 6 in the morning tomorrow.

Submit



156

Classify the following into fruits or vegetables, in the format specified below. 

Apple, Onion, Grapes, Tomato, Aubergine, Potato, Carrot

Format:

Banana <is a> fruit

Cabbage <is a> vegetable

Strawberries <are> fruit

Classification:

Submit Ctrl Enter

Submit



68

LLMs in practice

LLMs can perform all NLP “tasks”

And then some more



Information Retrieval

Ranking a list of documents or search results in response to an input query



Reading comprehension

Answering questions on a given text passage



Sentiment Analysis

Classifying the polarity or “happiness” of a given text



Summarization

Producing a shorter version of one or several documents that preserves most of the input’s meaning



Machine Translation

Translating a sentence in a source language to a different target language




Topic modelling

Discovering the abstract “topics” that occur in a collection of documents



Natural language inference

Determining whether a “hypothesis” is true (entailment), false (contradiction), or undetermined (neutral) given a “premise”




Relation prediction

Recognizing a named relation between two named semantic entities

Generative LLM use cases

Once you have a strong language model, the number of use cases is endless



Generative LLM use cases

Once you have a strong language model, the number of use cases is endless



Generative LLM use cases

Once you have a strong language model, the number of use cases is endless



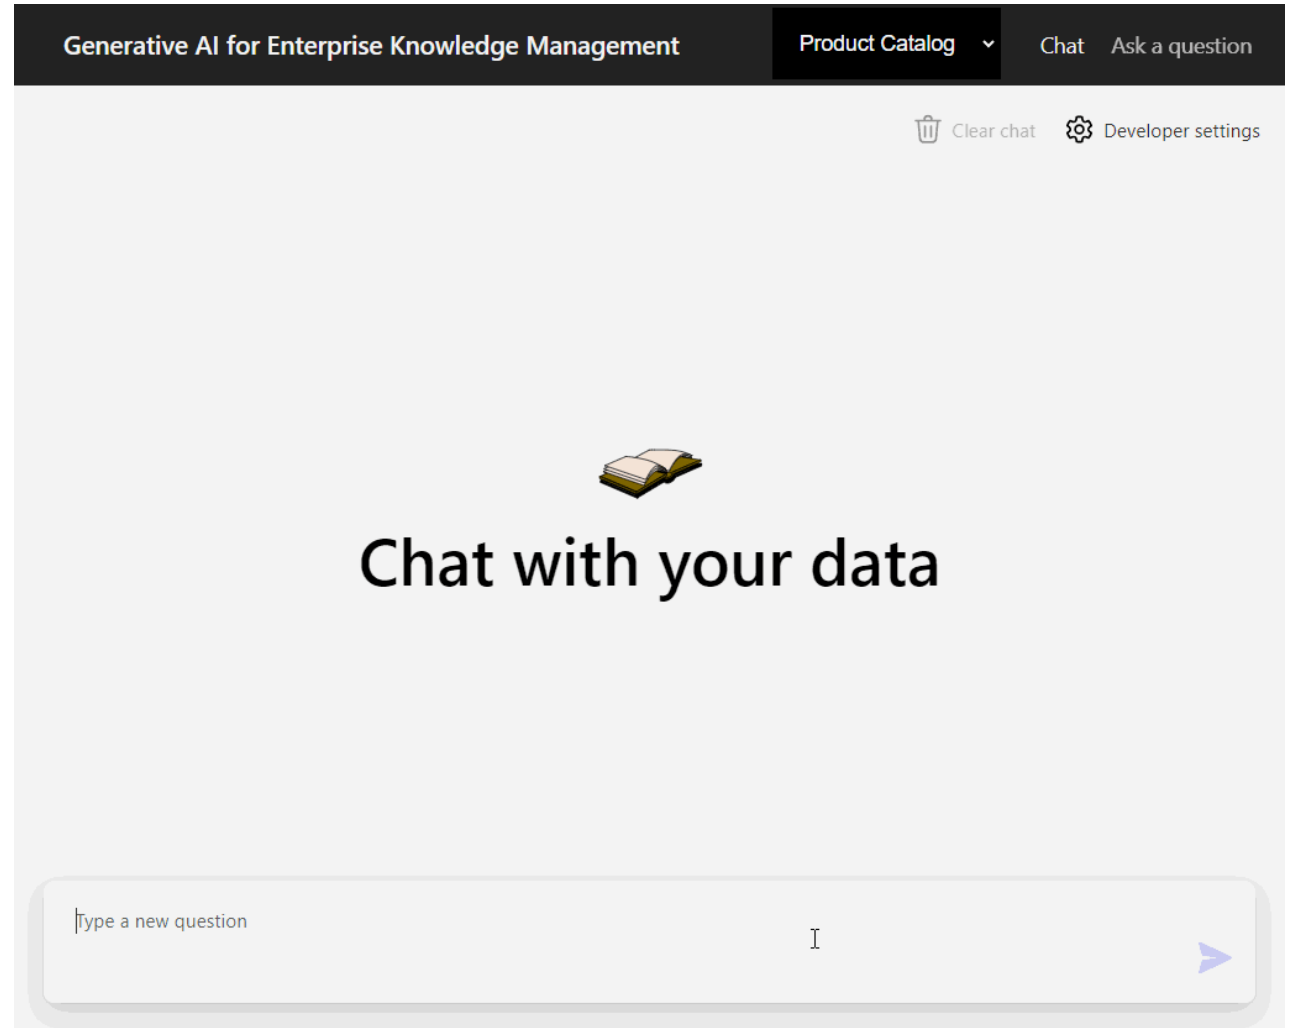
WRITING CODE



ENHANCED ON-PAGE SEARCH



**FINDING INTERNAL
INFORMATION**



Risks and limitations

Limitations

There are some limitations to consider when using LLMs

The cost per use of sophisticated models is materially significant. Fine-tuning the biggest models is even more expensive, potentially costing thousands of Euros a month.

Cost



Is the AI being used in a manner consistent with the purpose of the overall exercise? For e.g., submitting an AI-generated essay as your own.

Ethical Use



Models might output statements with confidence that are factually false. Sources and citations are unavailable for most models.

Hallucination and Confabulation



LLMs are comprised of billions of parameters. In theory, the larger the model, the better the output, but the compute time also increases.

Latency



Consent for data used (confidential information, personally identifiable information) is necessary, and the residency of data in geo-locations should also be compliant with legal and contractual requirements

Confidentiality & Privacy



SaaS-AI companies may use prompt payloads to train future versions of the base model, potentially including confidential data that could expose the user to IP infringement.

IP Protection & Infringement



A.I. TURNS THIS SINGLE
BULLET POINT INTO A
LONG EMAIL I CAN
PRETEND I WROTE.



A.I. MAKES A SINGLE
BULLET POINT OUT OF
THIS LONG EMAIL I CAN
PRETEND I READ.



TOM
FISH
BURNE

Tips and tricks

Some tips and tricks when applying NLP in practice

01

Dutch is more and more supported

Although English is best supported by NLP models, more and more multi-lingual and Dutch-based models are showing up, while translations are getting better as well

02

Use available libraries

Do not reinvent the wheel, reuse available libraries such as SpaCy, NLTK and LangChain

03

Use existing pipelines

Use the existing pipelines in e.g. Spacy and LangChain to deliver code that is easy to maintain and reuse

04

Sufficient data required

Consider if there's enough data to build a machine learning or deep learning model (rule of thumb: RB: hundreds, ML: thousands, DL: tens of thousands of labels, LLM: it depends)

05

Develop in an agile way

Start with the simplest approach that could possibly work (e.g. rule-based) and iterate on bigger and bigger models